

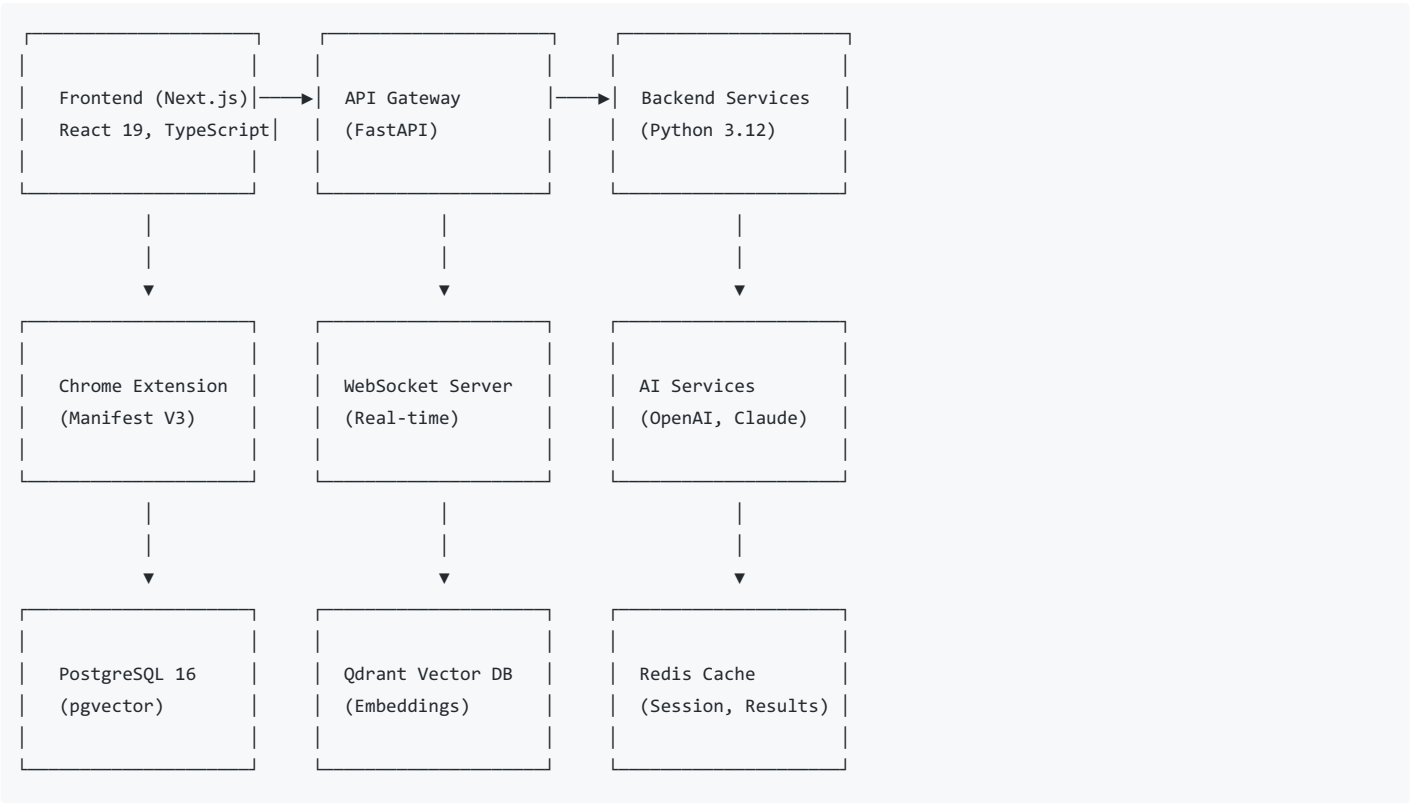
Promptitude Technical Architecture Overview

Table of Contents

- 1. [System Architecture](#)
- 2. [Core Features](#)
- 3. [Technology Stack](#)
- 4. [Data Flow](#)
- 5. [Performance Metrics](#)
- 6. [Security Architecture](#)

System Architecture

Promptitude is built on a modern microservices architecture designed for scalability, performance, and AI-first experiences.



Core Features

1. Mind Reader Search

What it does: Understands natural language queries and intent beyond keywords

- **Example:** "rockstar developer" → finds high-performing developers
- **Technology:** OpenAI embeddings + semantic analysis
- **Performance:** <300ms for 1M resumes

2. Progressive 3-Stage Search

What it does: Delivers results incrementally for optimal UX

- **Stage 1:** Instant results (<50ms) - cache hits
- **Stage 2:** Enhanced results (<200ms) - hybrid search
- **Stage 3:** Intelligent results (<500ms) - AI analysis
- **Technology:** Server-Sent Events (SSE) + async processing

3. Smart Talent Radar

What it does: Visual representation of candidate landscape

- **Visualization:** Interactive canvas-based radar
- **Interactions:** Zoom, rotate, click for details
- **Data Points:** Position (relevance), Color (availability), Size (experience)

4. Career DNA Matching

What it does: Finds candidates with similar career trajectories

- **Patterns:** Fast-track, Specialist, Lateral Explorer, Startup Builder
- **Algorithm:** 14-dimensional vector comparison
- **Accuracy:** 85% similarity matching

5. AI Interview Copilot

What it does: Real-time AI assistant during interviews

- **Features:** Live transcription, smart questions, fact-checking
- **AI Model:** GPT-4o-mini for real-time analysis
- **Latency:** <2 seconds for insights

6. Candidate Analytics

What it does: Advanced scoring beyond basic matching

- **Availability Score:** 0-1 scale based on profile indicators
- **Learning Velocity:** Speed of skill acquisition
- **Career Trajectory:** Pattern analysis and predictions

7. Chrome Extension

What it does: One-click LinkedIn profile import

- **Features:** Auto-detection, data extraction, duplicate prevention
- **Compatibility:** LinkedIn 2025 UI
- **Performance:** <3 seconds per import

Technology Stack

Frontend

- **Framework:** Next.js 15 with App Router
- **UI Library:** React 19 with Server Components
- **Language:** TypeScript 5.0
- **Styling:** Tailwind CSS 3.0
- **State Management:** Zustand
- **Real-time:** Socket.IO client

Backend

- **Framework:** FastAPI (Python 3.12)
- **ORM:** SQLAlchemy 2.0 with async support
- **Task Queue:** Celery with Redis
- **API Design:** RESTful with OpenAPI 3.0
- **Authentication:** JWT with refresh tokens

Databases

- **Primary:** PostgreSQL 16 with pgvector extension
- **Vector Store:** Qdrant (1536-dim embeddings)
- **Cache:** Redis 7.0
- **Search Index:** PostgreSQL full-text search

AI/ML Stack

- **Embeddings:** OpenAI text-embedding-ada-002
- **Chat Models:** GPT-4o-mini, Claude 3.5 Sonnet
- **Frameworks:** LangChain, OpenAI Python SDK
- **Monitoring:** Weights & Biases

Infrastructure

- **Deployment:** Docker containers on Railway
- **CDN:** Cloudflare
- **Monitoring:** Datadog, Sentry
- **CI/CD:** GitHub Actions

Data Flow

Search Flow

User Query → Query Parser → Skill Extraction → Parallel Processing:

- └─ Cache Check (Redis)
- └─ Keyword Search (PostgreSQL BM25)
- └─ Vector Search (Qdrant)

↓

Result Fusion → Scoring → Analytics Enhancement → AI Enhancement → Response

Import Flow

LinkedIn Profile → Chrome Extension → Data Extraction → Validation →
Backend API → Duplicate Check → Data Cleaning →
Database Insert → Vector Embedding → Qdrant Index → Success Response

Interview Flow

Audio Stream → WebSocket → Transcription Service →
Text Buffer → AI Analysis (GPT-4) →
Insight Generation → WebSocket → UI Update

Performance Metrics

Search Performance

- **Instant Stage:** <50ms (p99)
- **Enhanced Stage:** <200ms (p99)
- **Intelligent Stage:** <500ms (p99)
- **Throughput:** 1000 searches/second
- **Concurrent Users:** 10,000

Vector Search

- **Index Size:** 1M vectors
- **Query Time:** <100ms
- **Accuracy:** 95% relevance
- **Memory Usage:** 8GB for 1M vectors

Database Performance

- **Write Speed:** 5000 resumes/minute
- **Query Response:** <10ms for indexed queries
- **Connection Pool:** 100 connections
- **Cache Hit Rate:** 60%

Security Architecture

Authentication & Authorization

- **Method:** JWT tokens with refresh mechanism
- **Token Lifetime:** 15 minutes (access), 7 days (refresh)
- **Permissions:** Role-based access control (RBAC)
- **MFA:** Optional TOTP support

Data Protection

- **Encryption at Rest:** AES-256
- **Encryption in Transit:** TLS 1.3
- **PII Handling:** Tokenization for sensitive data
- **GDPR Compliance:** Right to deletion, data portability

API Security

- **Rate Limiting:** 100 requests/minute per user
- **Input Validation:** Pydantic models
- **SQL Injection:** Parameterized queries
- **XSS Prevention:** Content Security Policy

Chrome Extension Security

- **Permissions:** Minimal required (activeTab, storage)
- **Content Security:** Isolated content scripts
- **Data Transit:** Encrypted API calls
- **Auth Storage:** Secure chrome.storage

Scalability Considerations

Horizontal Scaling

- **API Servers:** Auto-scaling with load balancer
- **Vector Search:** Distributed Qdrant cluster
- **Database:** Read replicas for search queries
- **Cache:** Redis cluster with sharding

Performance Optimization

- **Query Caching:** 1-hour TTL for common searches
- **Connection Pooling:** Persistent DB connections
- **Lazy Loading:** Progressive data fetching
- **CDN:** Static assets and API responses

Future Enhancements

1. **Multi-region Deployment:** Reduce latency globally
2. **GraphQL API:** Efficient data fetching
3. **ML Model Caching:** Edge deployment for embeddings
4. **Real-time Sync:** WebRTC for interview features
5. **Blockchain Integration:** Verified credentials

Monitoring & Observability

Metrics Collection

- **APM:** Datadog for application performance
- **Logs:** Centralized with ELK stack
- **Errors:** Sentry for exception tracking
- **Custom Metrics:** StatsD for business metrics

Alerting

- **Uptime:** 99.9% SLA monitoring
- **Performance:** Latency and error rate alerts
- **Security:** Anomaly detection for auth failures
- **Business:** Search quality degradation alerts

Development Workflow

Local Development

```
# Backend
cd backend
python -m venv venv
source venv/bin/activate
pip install -r requirements.txt
uvicorn app.main:app --reload

# Frontend
cd frontend
npm install
npm run dev

# Vector Database
docker run -p 6333:6333 qdrant/qdrant
```

Testing Strategy

- **Unit Tests:** 85% coverage target
- **Integration Tests:** API endpoint testing
- **E2E Tests:** Playwright for critical paths
- **Load Tests:** K6 for performance testing

This technical overview provides the foundation for understanding Promtitude's architecture. For detailed information about specific features, refer to the individual documentation files in this directory.